

Seminar : Concurrency Control

There are several problems that ^{are} found when concurrent transaction execute in an uncontrolled manner.

Let's take an example of a simplified air line reservations database in which a record is stored for each airline flight.

* Each record includes the No of Reserved seats on that flight^{assigned} (uniquely identifiable) data item, among other information

T₁

read-item(X)

$X = X - N$

write-item(X)

read-item(Y)

$Y = Y + N$

write-item(Y)

T₁ transaction that transfers N reservations from one flight whose No of Reserved seats is stored in db item named X to another flight whose no of Reserved seats is stored in db item named Y.

T₂

read-item(X)

$X = X + M$

write-item(X)

T₂ transaction that just reserves M seats on the first flight (X)

So when a db access program is written it has a flight number, flight date and no of seats to be booked as parameters

hence the same program can be used to execute many different transactions, each with a different flight number, flight date & No of seats to be booked

For concurrency control purpose, a transaction is a particular execution of a program on a specific date, flight and number of seats.

T₁ and T₂ are specific execution of the program that refer to specific flights whose no of seats are stored in data item X and Y.

When these 2 simple transaction execute concurrently, we may encounter many problems

1. The Lost Update Problem:-

When 2 transaction that access the same db items have their operations interleaved such that ^{it makes} value of some db items incorrect.

If $X = 80$ at start

$N = 5$

$X = 80 - 5 = 75$

$M = 4$

$X = 75 + 4 = 79$

should be the final result

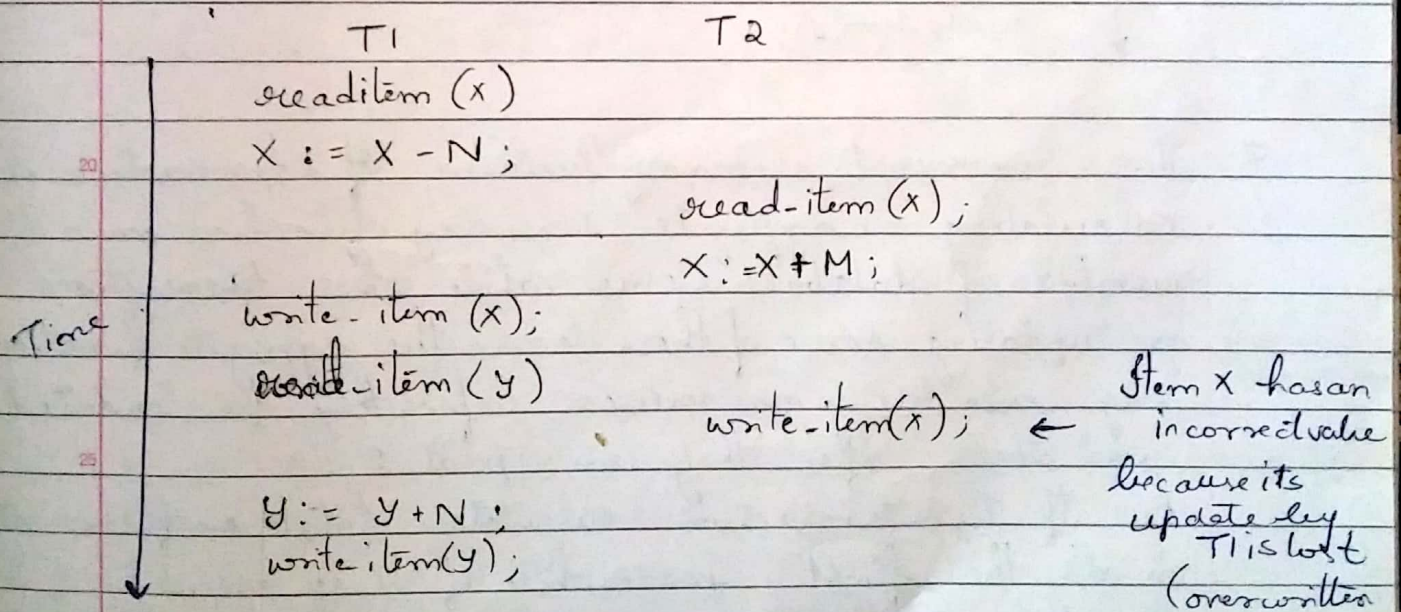
However in the interleaving of operation in T_1 and T_2

$X = 80$

$X = 80 + 4 = 84$ coz the

update in T_1 that removed 5 seats was lost.

Suppose transactions T_1, T_2 submitted at approximately same time and suppose operations are interleaved, the final value of X is incorrect, as T_2 reads the value of X before T_1 changes in db, hence updated value resulting from T_1 is lost.



2. The Temporary Update (or Dirty Read) Problem.

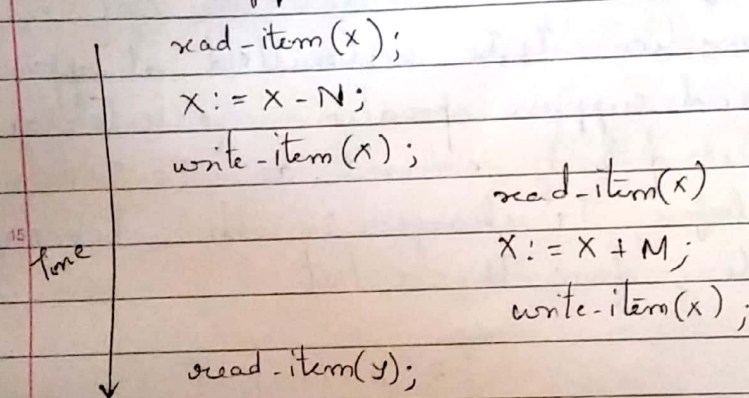
When 1 transaction updates a db item and then transaction fails for some reason.

Meanwhile, the updated item is accessed (read) by another transaction before it is changed/rolled back to its original value.

This value of db item is called as dirty data since it has been created by a transaction that has ^{not} completed & committed yet. Hence the problem is referred as dirty

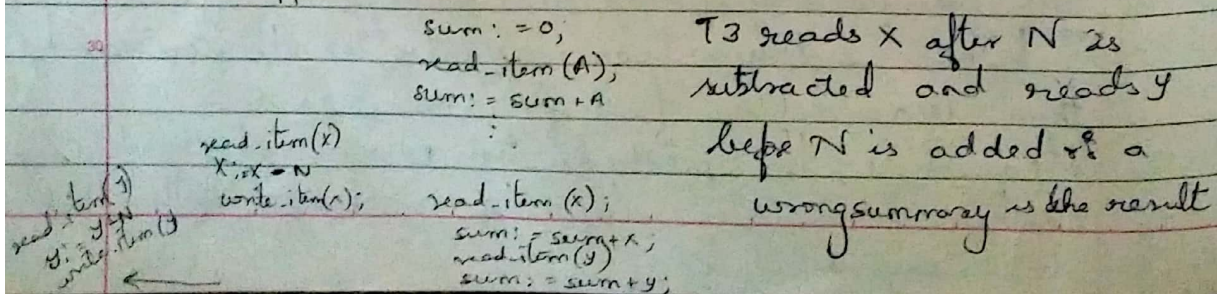
5 read problem

Eg: T_1 updates item X and fails before completion at such case, the system rolls back changes. But before it can do so, transaction T_2 reads temporary/updated value of X which will not be recorded permanently in db item as T_1 failed. The value of item X read by T_2 is called dirty data.



3. The Incorrect Summary Problem: If 1 transaction is calculating an aggregate summary function on a number of database items while other transactions are updating some of these items, the aggregate function may calculate some values before they are updated and others after they are updated.

Eg: If T_3 , transaction calculates total No of reservation on all the flight, meanwhile if T_1 is executing. If interleaving of operation occurs as shown, the result



The Unrepeatable Read Problem

If a Transaction T reads the same item twice & the item is changed by another transaction T' between the 2 reads. Hence T receives different values for its 2 reads at the same time.

Eg: If during an airline reservation transaction a customer inquires seat availability on different flights.

When customer decides on a particular flight, the transaction then reads the No of seats on flight at a time before completing reservation and it may end up reading a different value for the item.